The Contract: Defining Scope & Constraints

FUNCTIONAL REQUIREMENTS (The Product)

- **1:1 & Group Chat:** Support for multi-user participation (Cap: 100 users/group).
- **Message Delivery:** Text, Images, and Video support.
- **Offline Support:** Store & Forward mechanism for disconnected users.
- **Ack Mechanism:** Sent, Delivered, and Read receipts.

NON-FUNCTIONAL REQUIREMENTS (The Engineering)

- **Scale:** 1 Billion DAU. ~40k msg/sec average, ~230k+ peak RPS.
- **Latency:** Real-time delivery target <500ms (P99).
- **Reliability:** 'At least once' delivery guarantee. Zero message loss.
- **Ephemeral Storage:** Strict privacy policy. Messages deleted from server upon delivery (Max retention: 30 days).

The Protocol: Why We Don't Use REST

HTTP Polling (Inefficient)

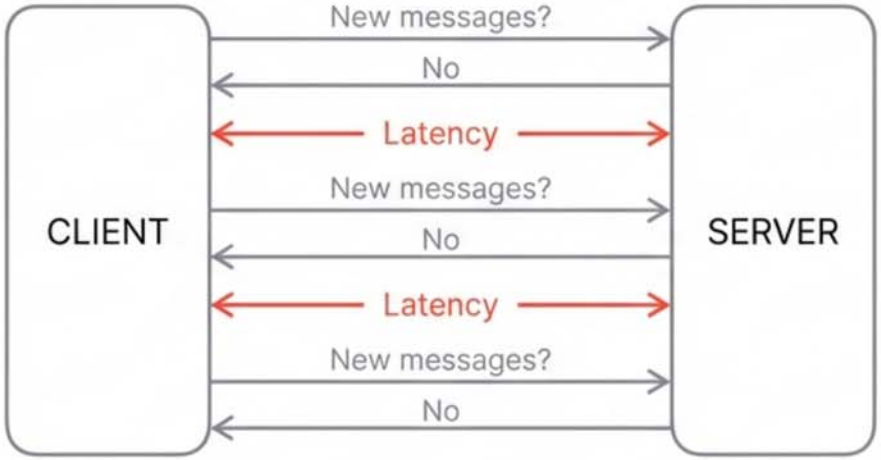

WebSocket (Persistent)

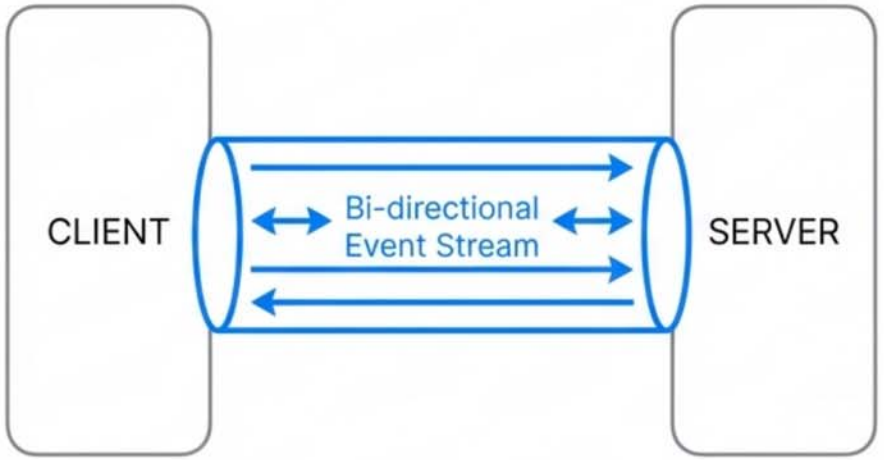

The API Contract (Events)

```
-> createChat(participants) : returns chatId
-> sendMessage(chatId, content) : returns status
<- receiveMessage(chatId, senderId, content) : [SERVER PUSH]
-> ackMessage(messageId) : [CLIENT CONFIRM]
```

The Data Model: Optimizing for Write Heavy Workloads

Database Choice: DynamoDB (Key-Value / Wide-Column)

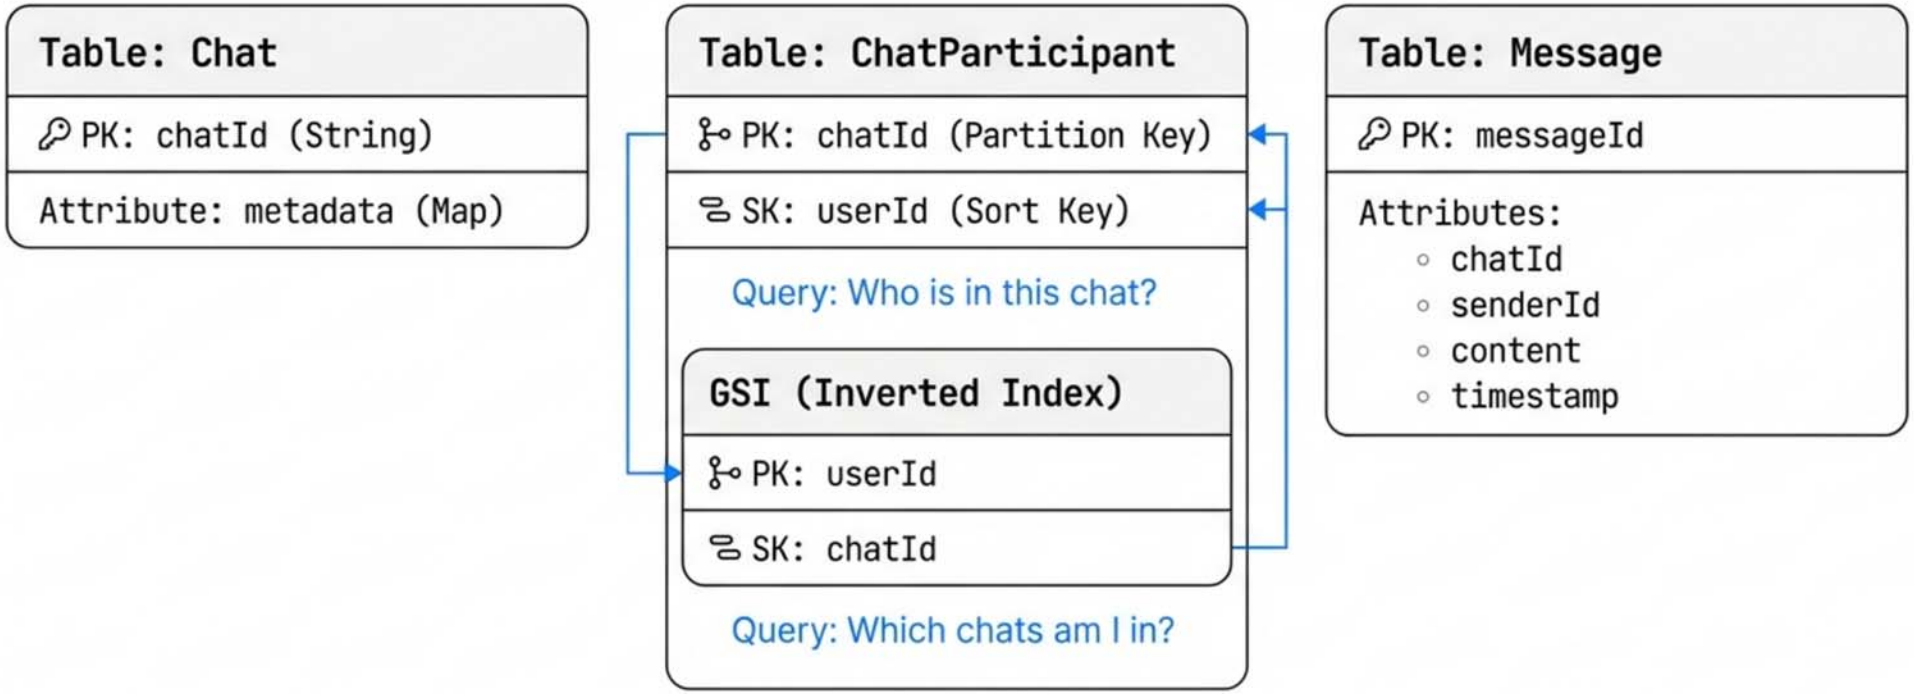

Reasoning: DynamoDB handles massive horizontal scaling for write-heavy chat logs better than traditional SQL joins.

The Functional Baseline (The MVP)

Act II: A Single Server Solution

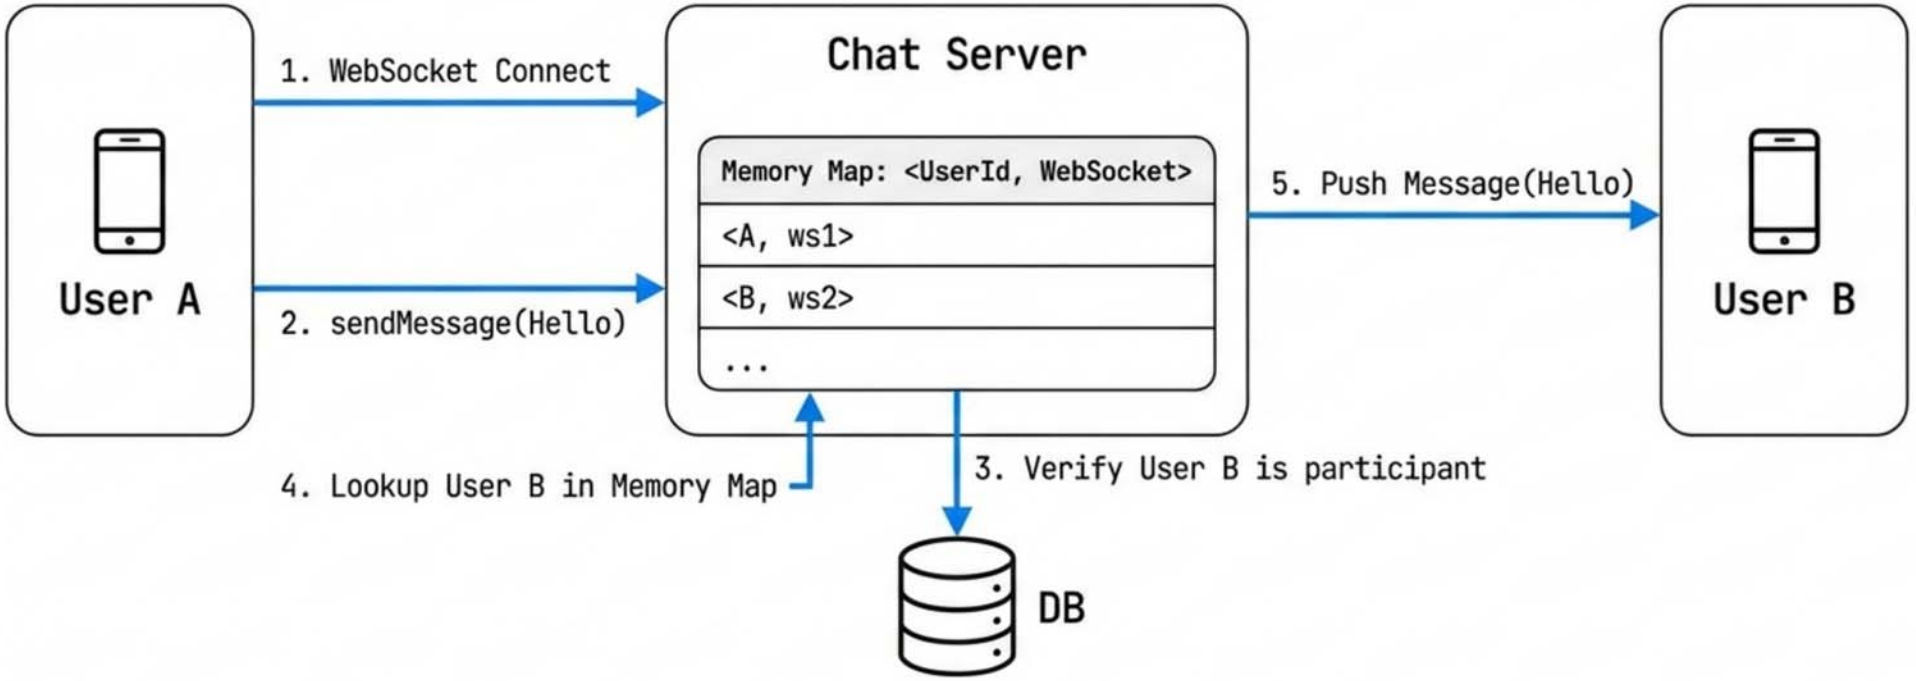

CONSTRAINT: Works for 500 users. Fails at 1 Billion (Memory/CPU limits).

The ‘Offline’ Problem: Store and Forward

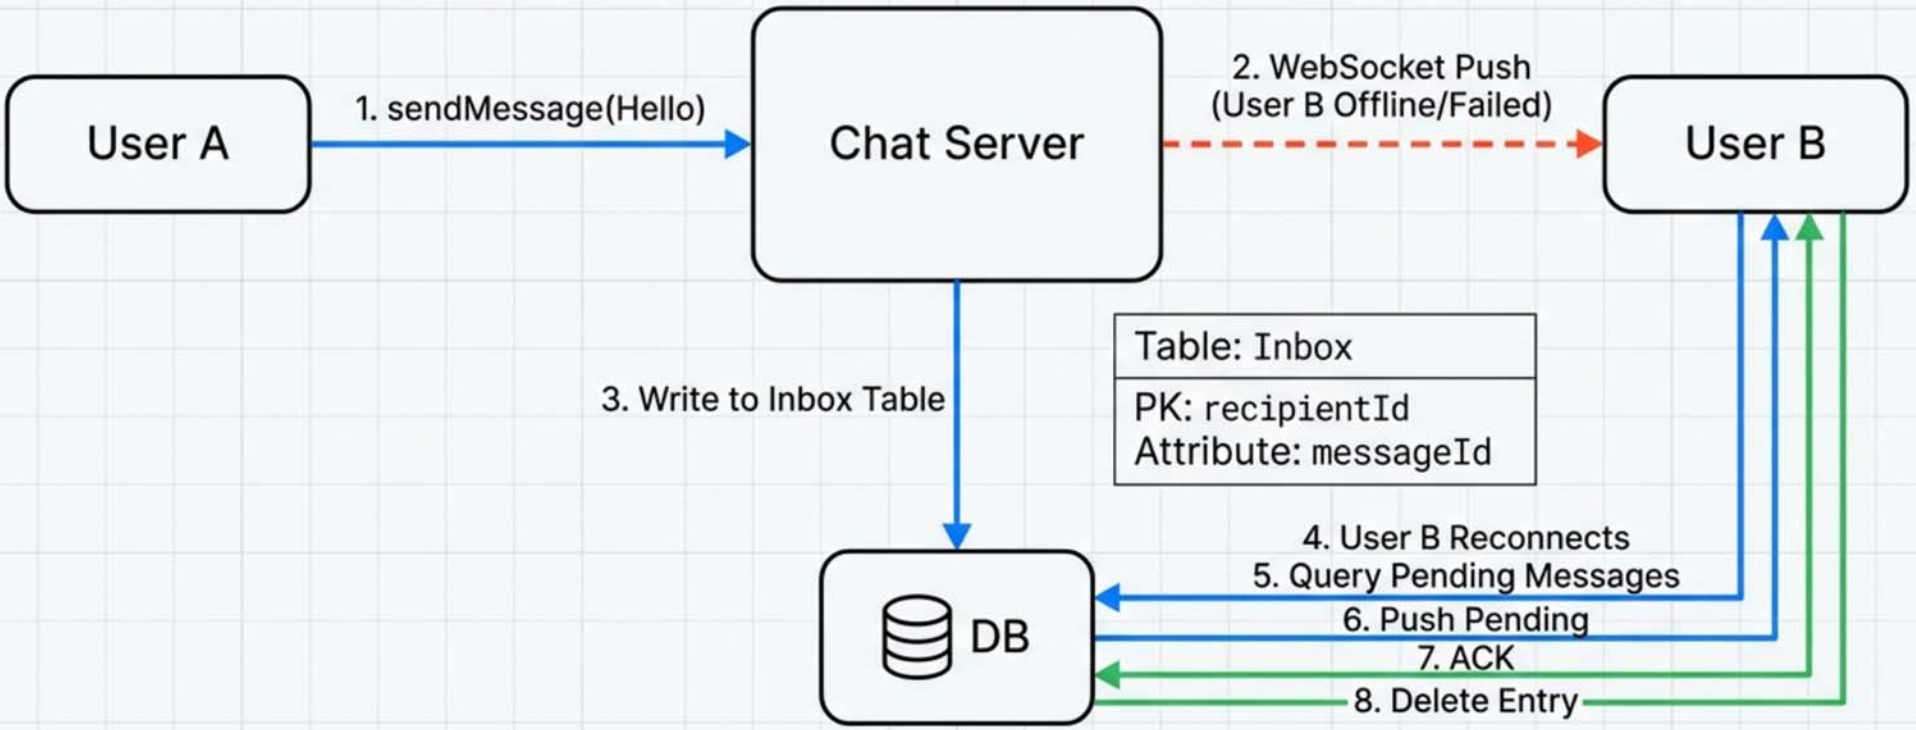

Handling Media: Decoupling Data from Signaling

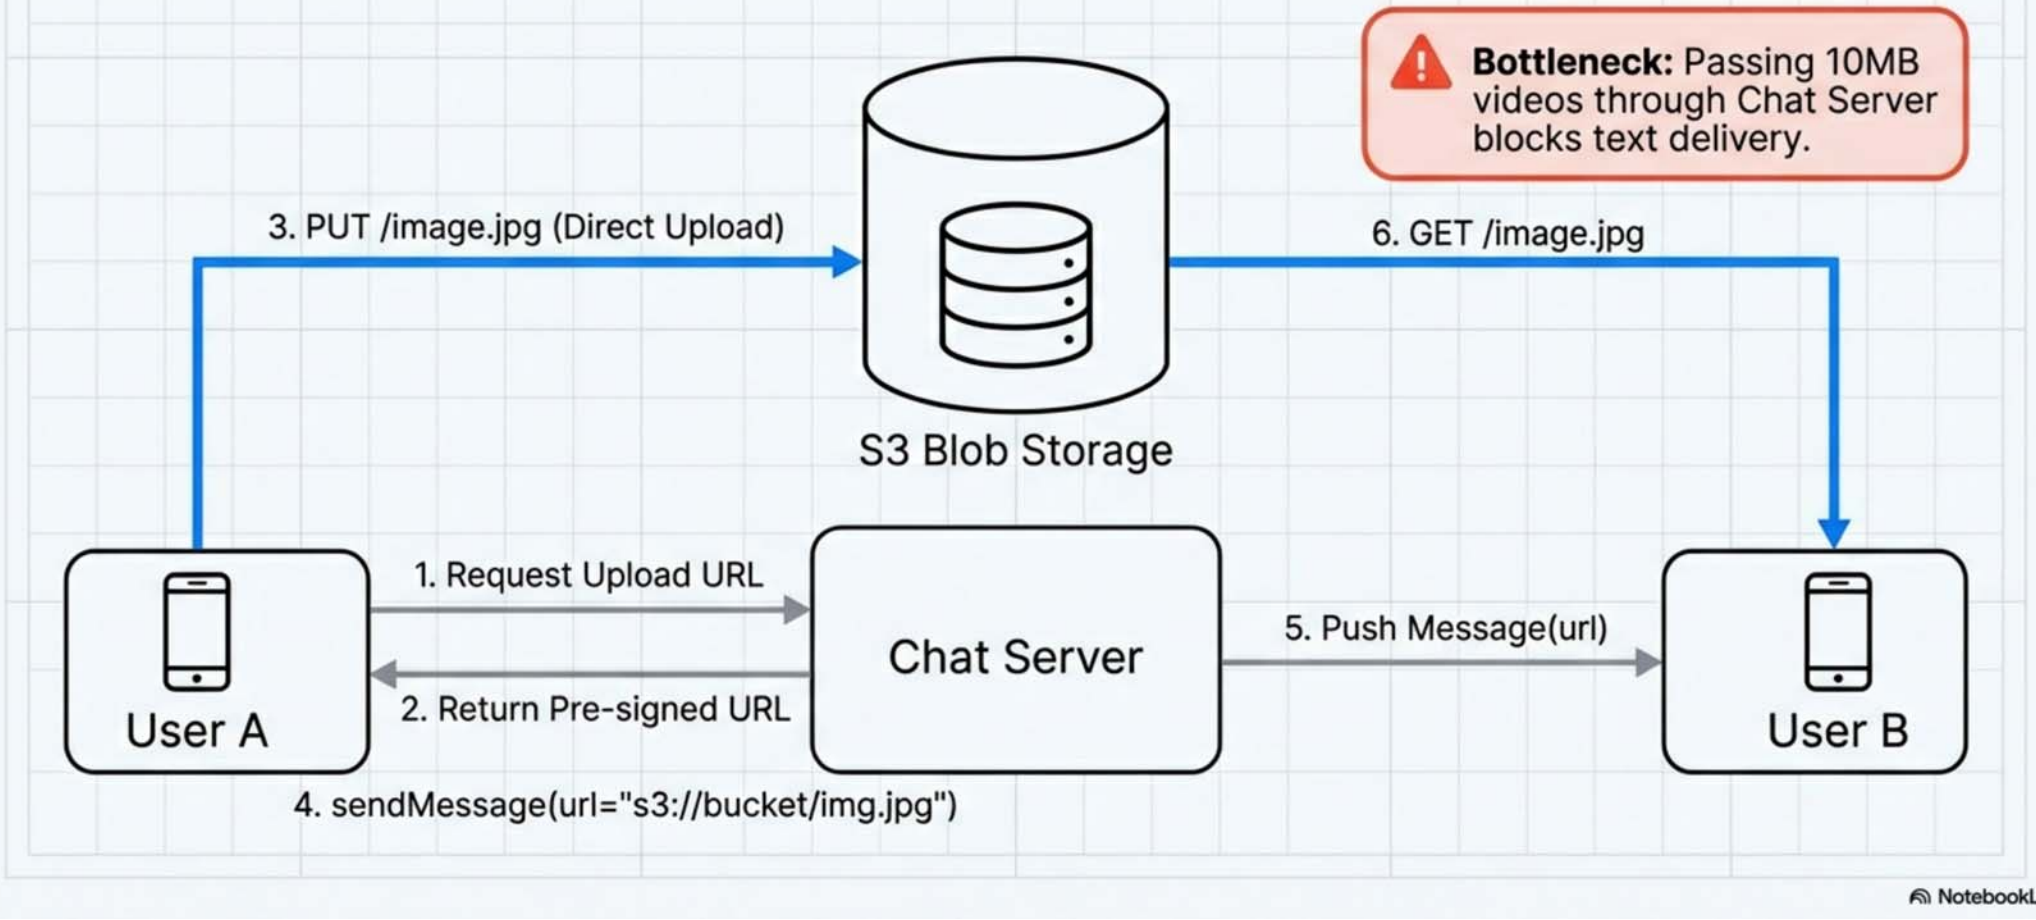

Bottleneck: The 1 Billion User Limit

Scaling Connections with Layer 4 Load Balancing

THE MATH:

1 Billion Users = ~100M Concurrent Connections

Single Server Cap = ~2M Connections

Requirement = ~50-100 Servers

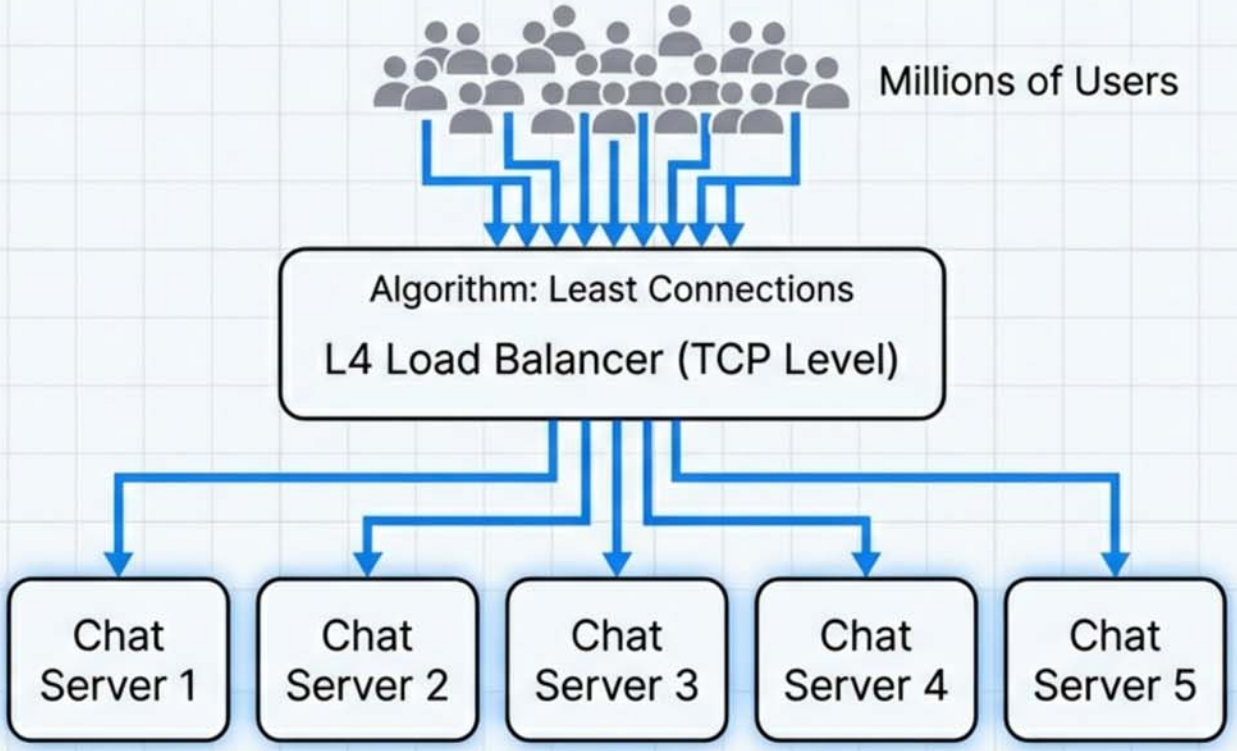

The Routing Challenge: Bridging the Islands

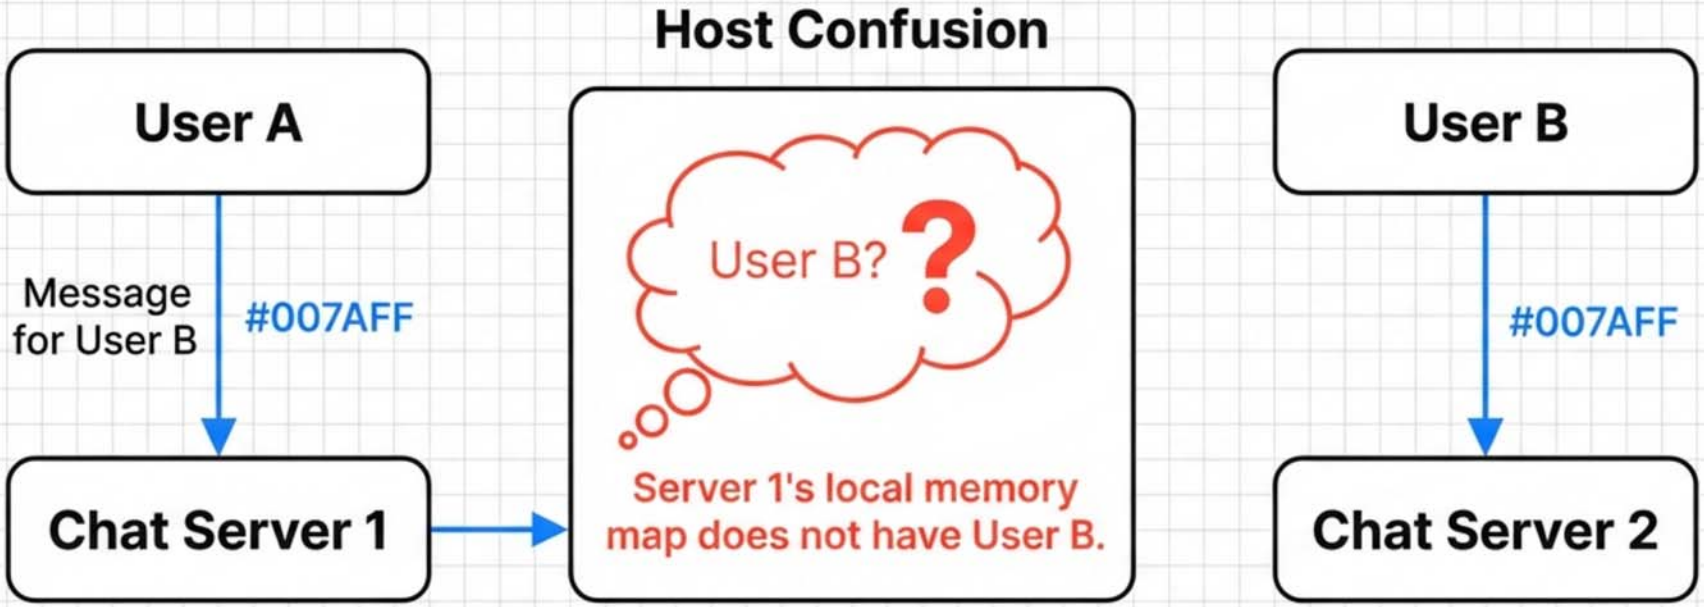

WHY NOT KAFKA? Kafka topics are heavy. Creating **`1 Billion`** topics (one per user) is inefficient. Kafka is for streams, not ephemeral routing.

The Solution: Redis Pub/Sub

Turning isolated servers into a unified mesh.

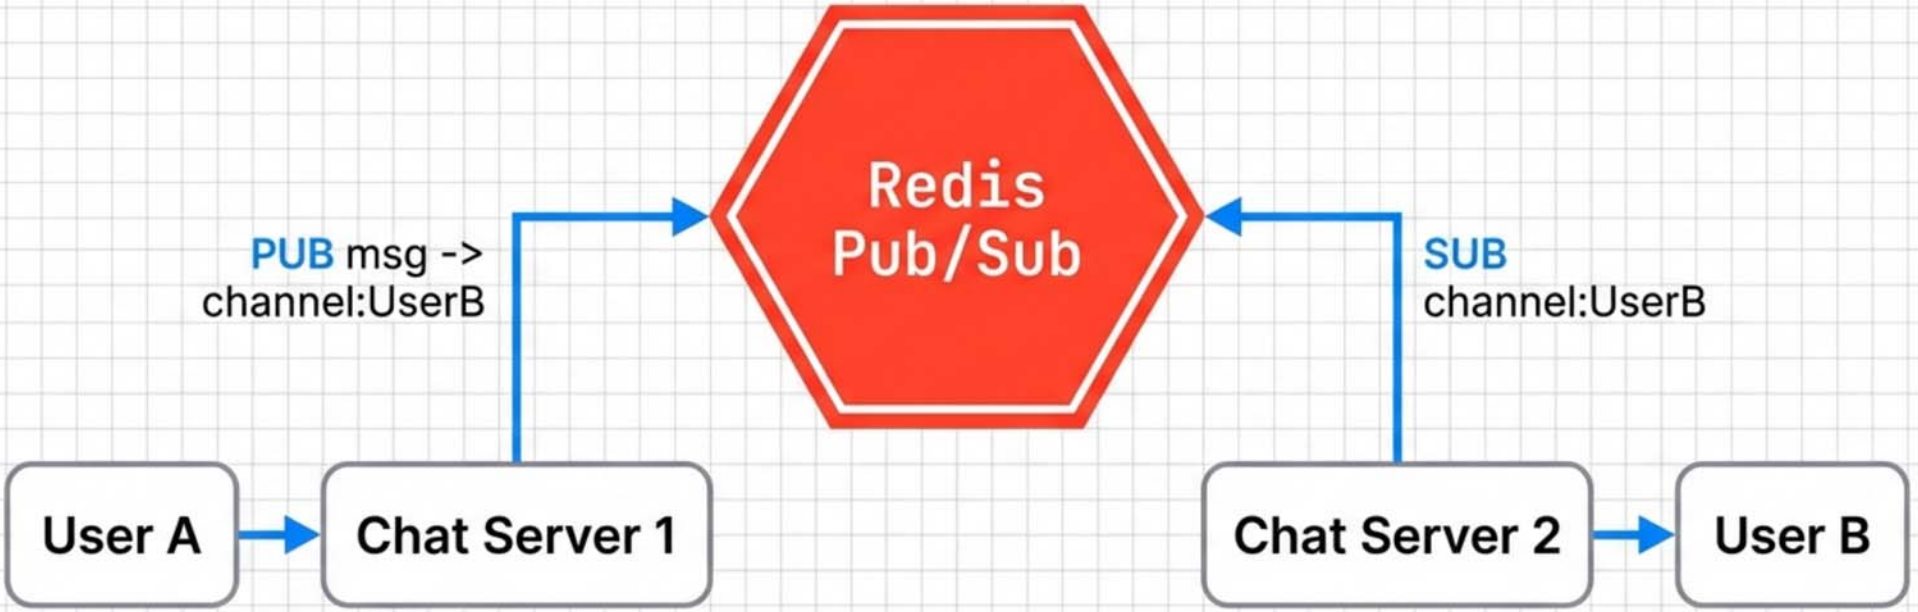

Redis acts as a lightweight, ephemeral message bus.

Scaling Redis: Sharding and Clusters

⚠️ Single Redis Node = Memory/CPU Bound.

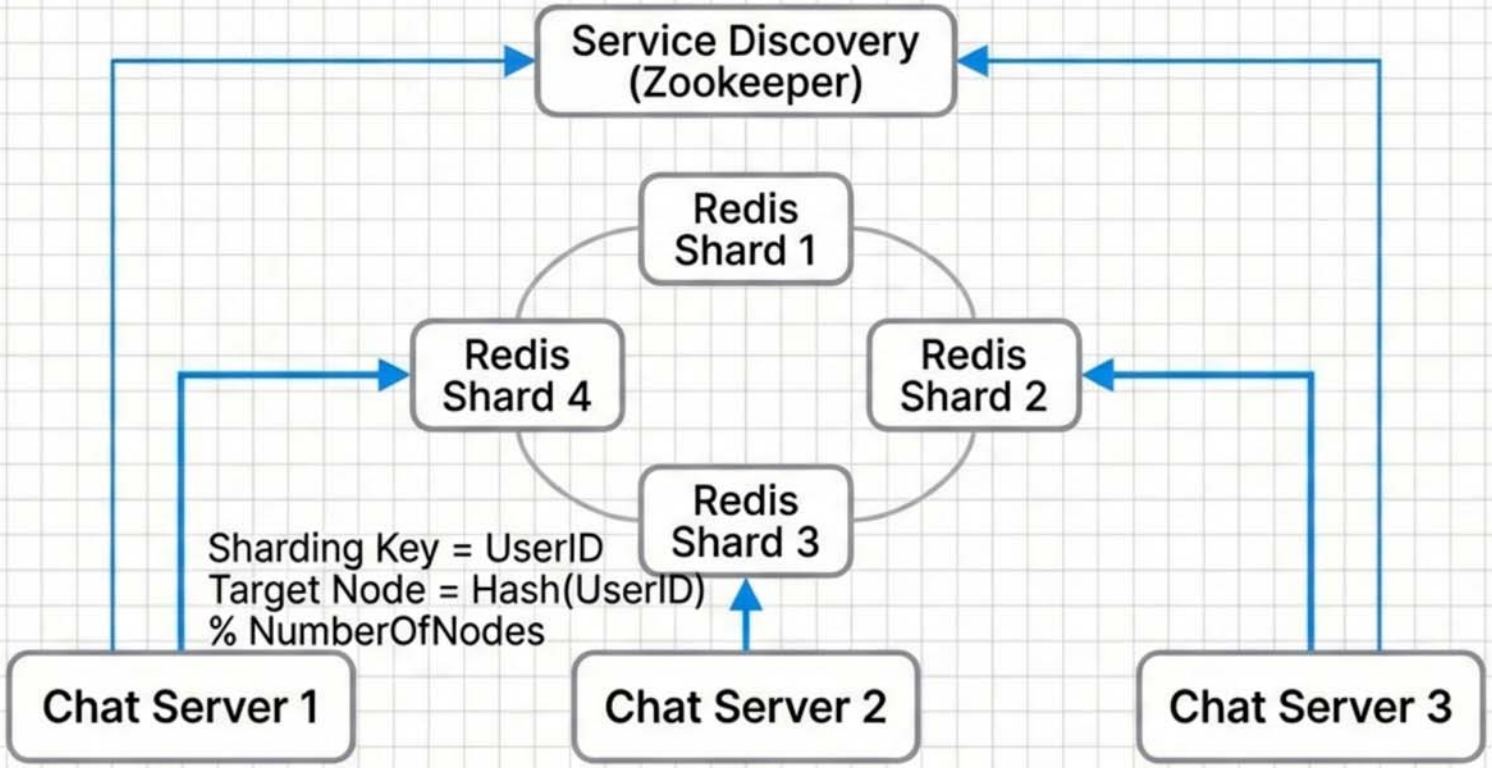

Connection Pooling enabled to minimize TCP handshake overhead.

Reliability: Solving “At Most Once” Delivery

Implementing a persistent messaging strategy to prevent data loss.

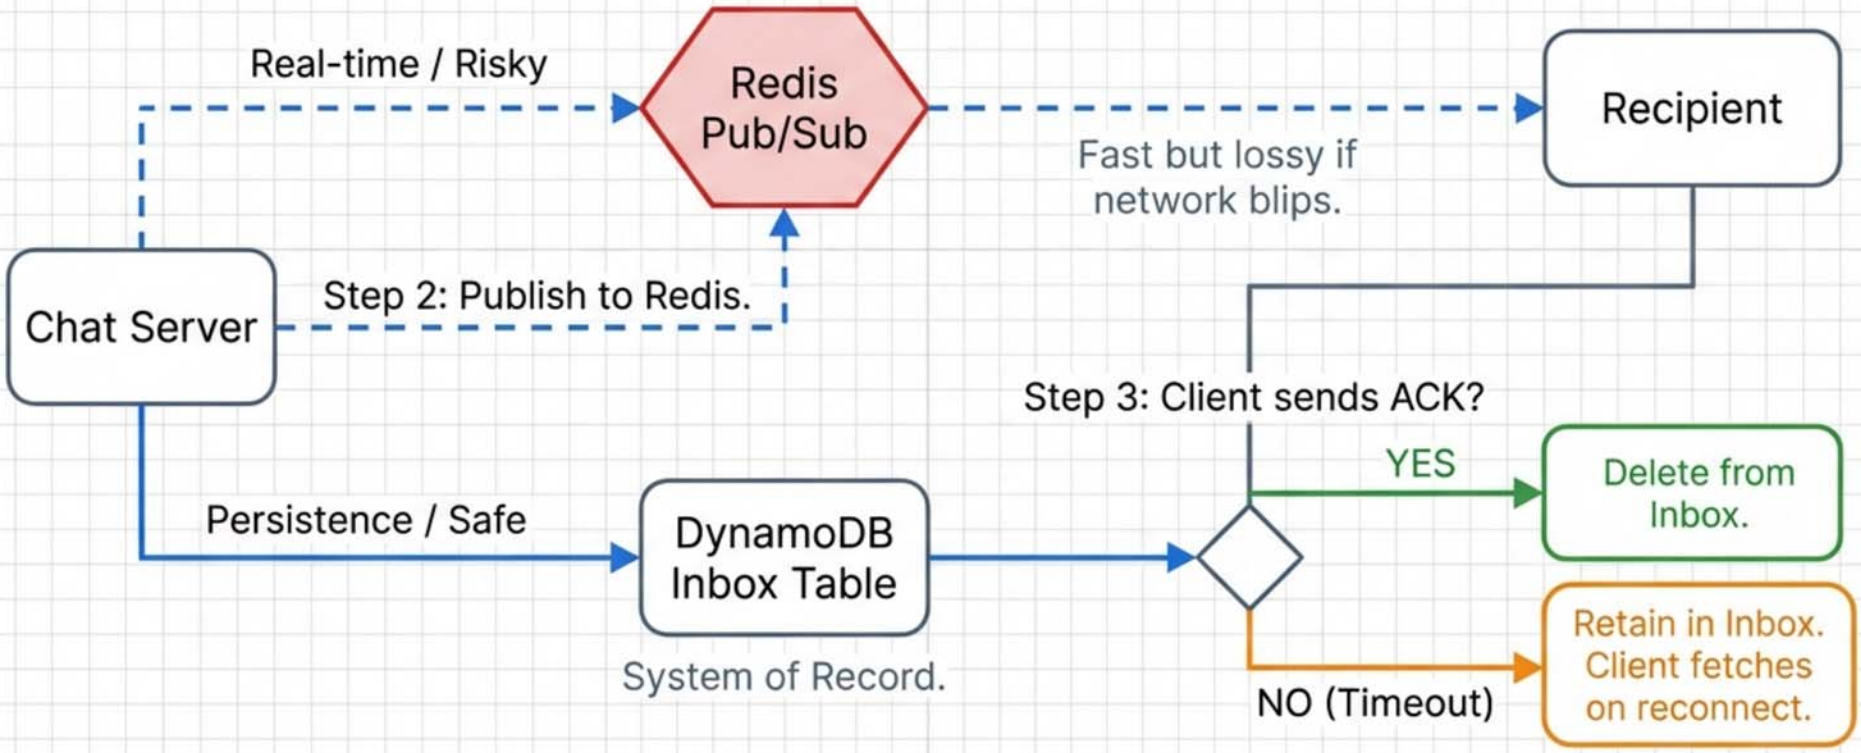

Multi-Device Support: Fan-out Complexity

Implementing a multi-dart messaging strategy to prevent data loss.

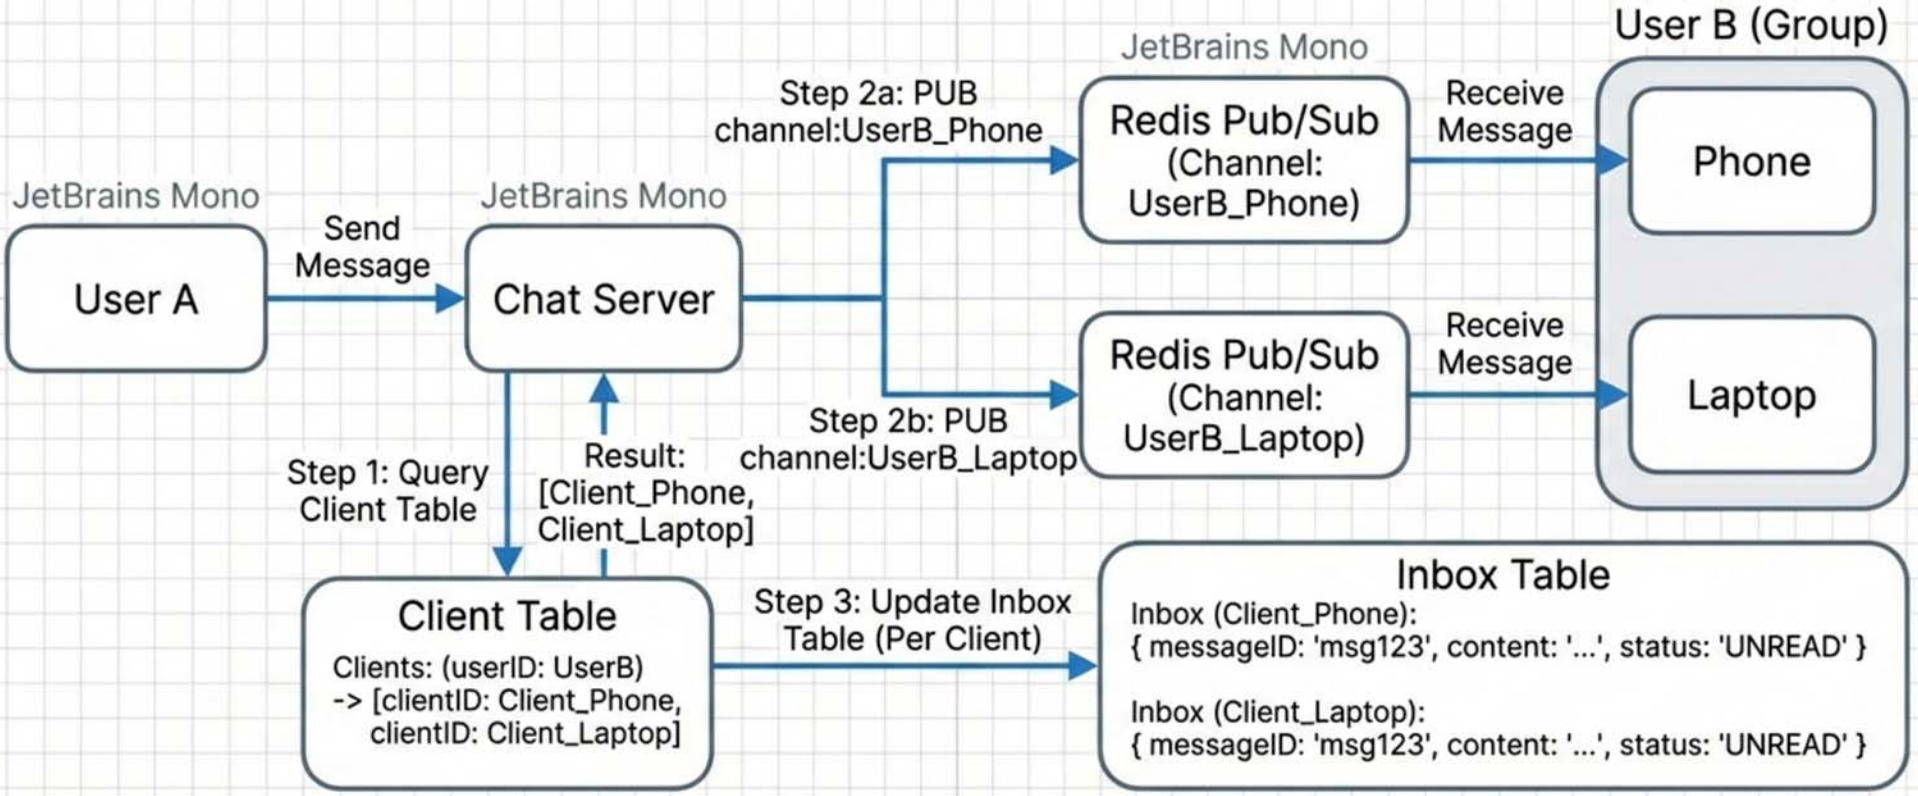

Ephemeral Storage & Cleanup

Privacy Principle: Data is Toxic Sludge.

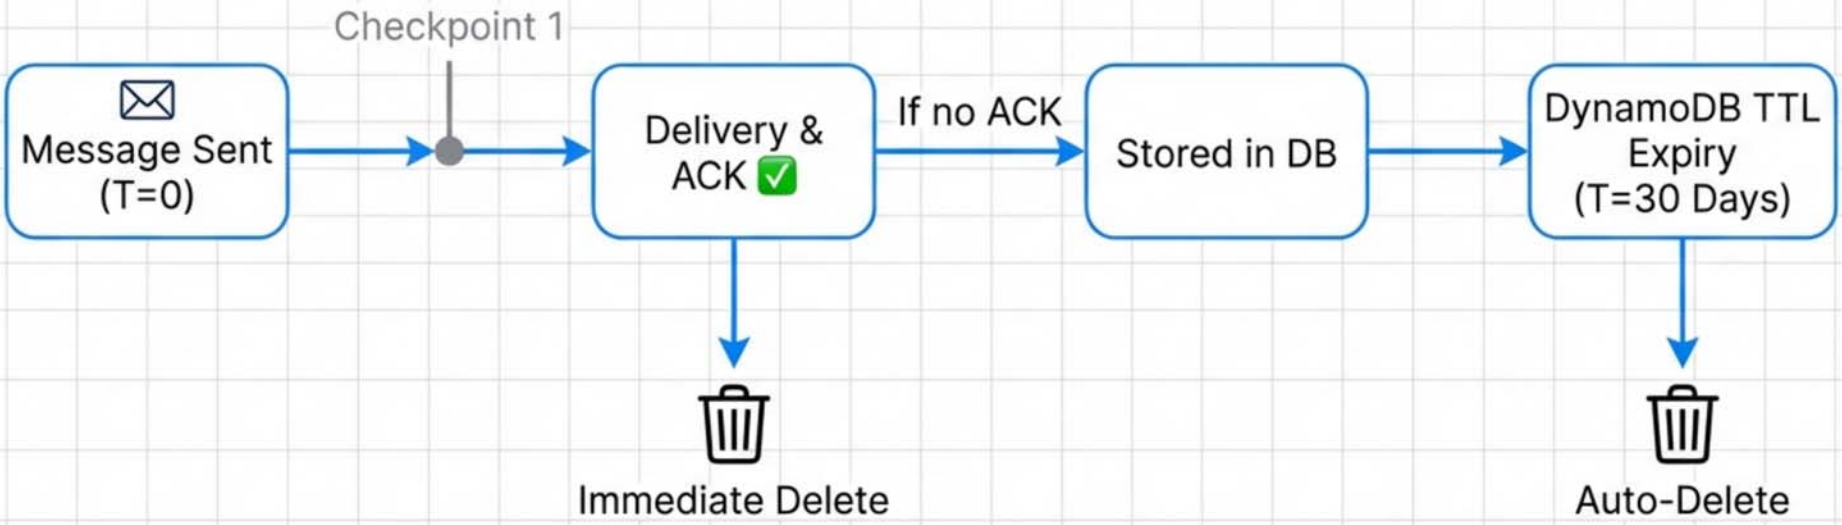

Mechanism: DynamoDB TTL (Time To Live). Database engine automatically expires rows based on timestamp. No custom cleanup code required.

The Final Architecture

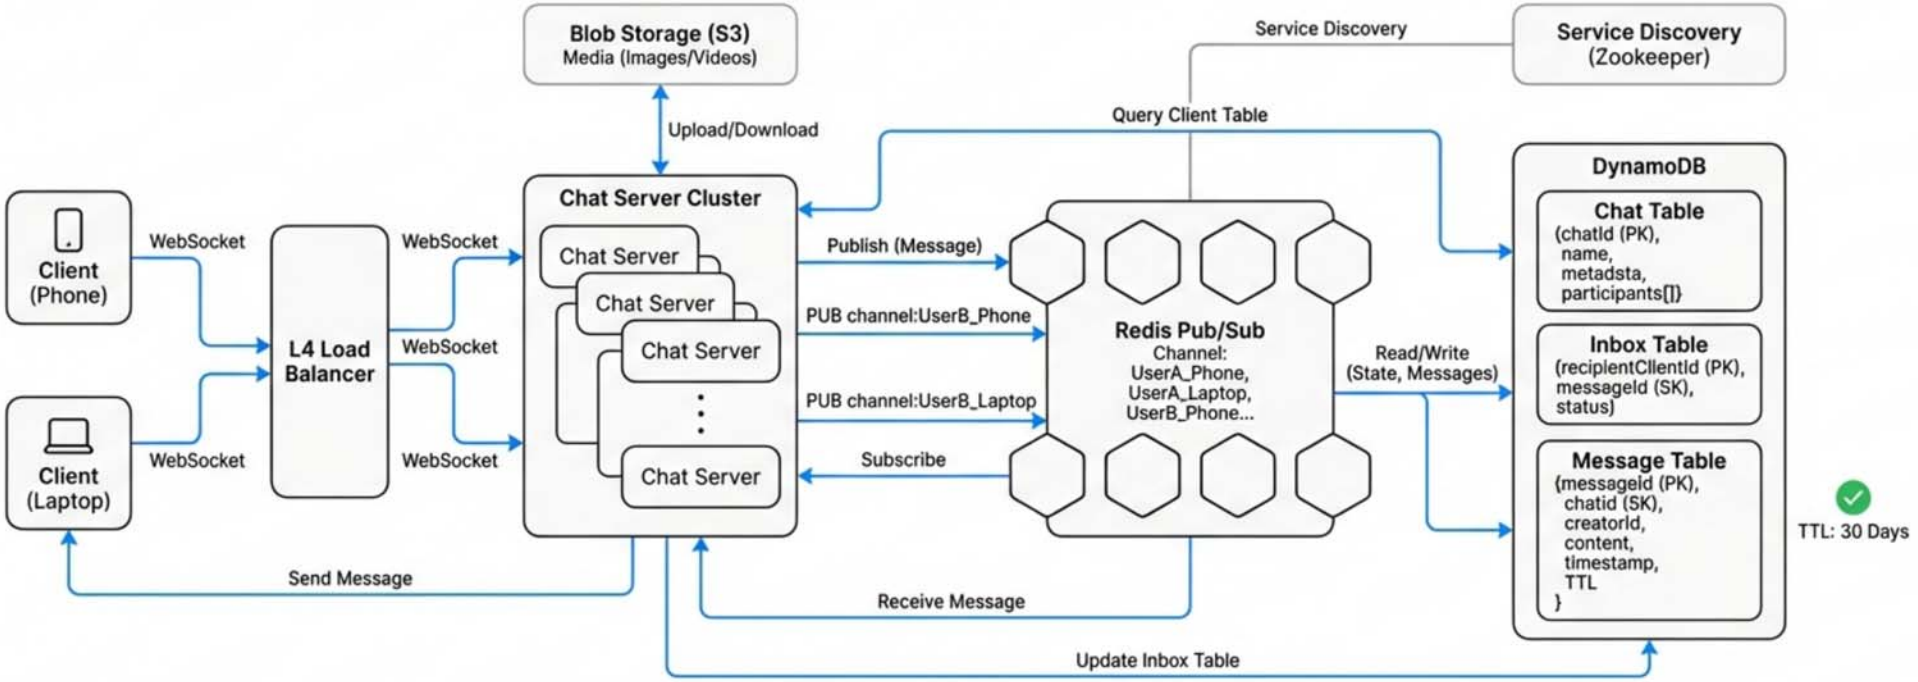

Summary & Trade-offs

PROS:
- Low latency (<500ms)
- Horizontal Scale (1B Users)
- High Reliability (Hybrid DB/Redis)

CONS:
- Operational complexity of Redis Cluster
- Client-side logic needed for message ordering