

PHASE 1: THE CONTRACT

Core Features (IN SCOPE)

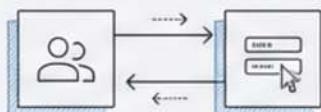
-  Shorten: Long URL → Short Alias
-  Redirect: Short Alias → Long URL
-  Custom Alias (Optional)
-  Expiration Date (Optional)

Out of Scope

-  User Authentication
-  Complex Analytics Dashboards
-  Link Editing

NON-FUNCTIONAL REQUIREMENTS

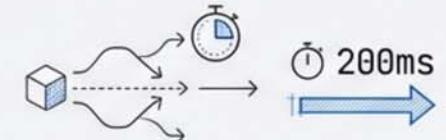
Scale: 100M Daily Active Users | 1 Billion URLs



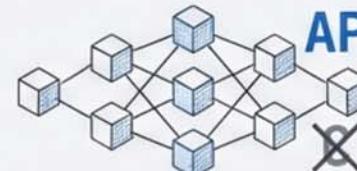
100:1

Read : Write Ratio

Latency: < 200ms (Redirects)

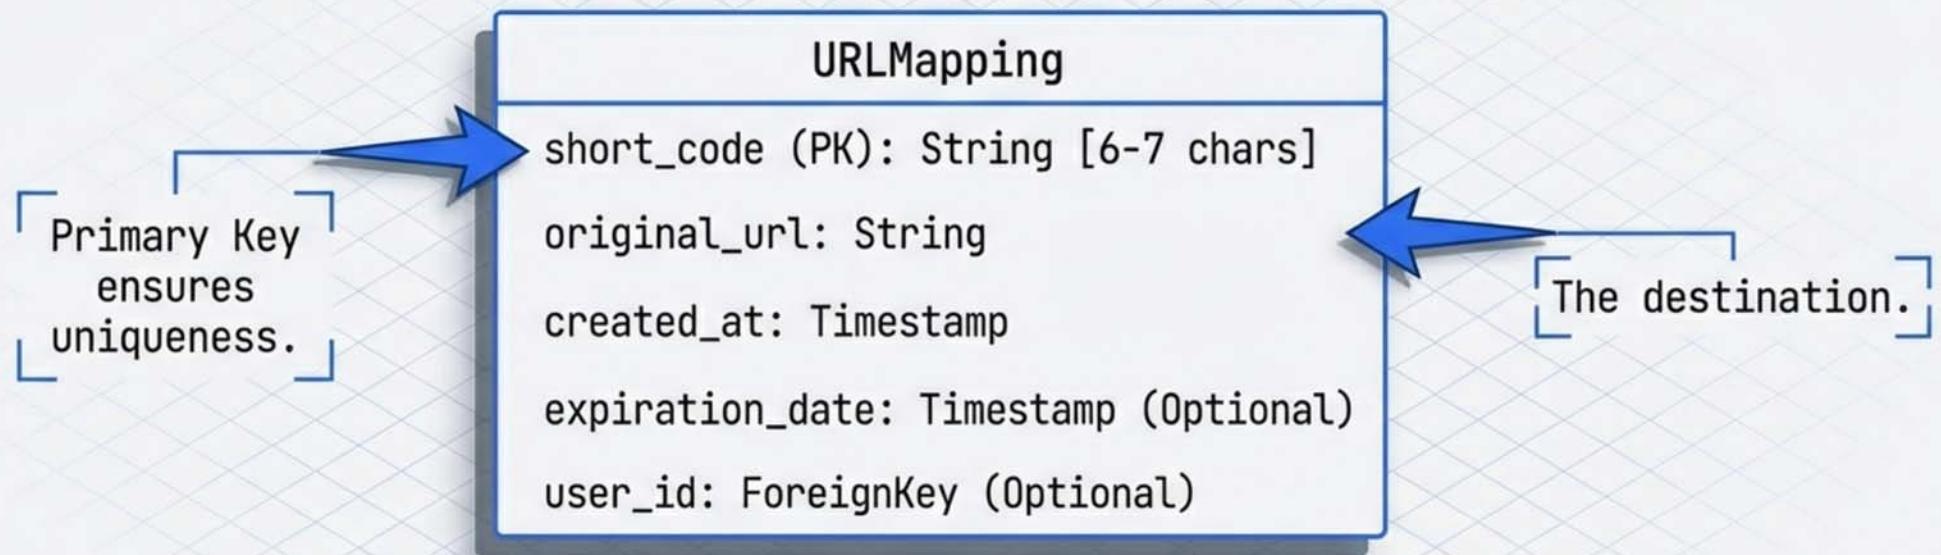


Availability: AP System
(High Availability > Consistency)



Read-Heavy System

Phase 2: The Core Entities



Phase 3: The API Contract

Create Short URL

```
POST /api/urls/shorten  
Request Body:  
{  
  "long_url": "https://www.example.com/....,  
  "alias": "custom-name" (optional)  
}  
Response: 200 OK  
{ "short_url": "bit.ly/custom-name" }
```

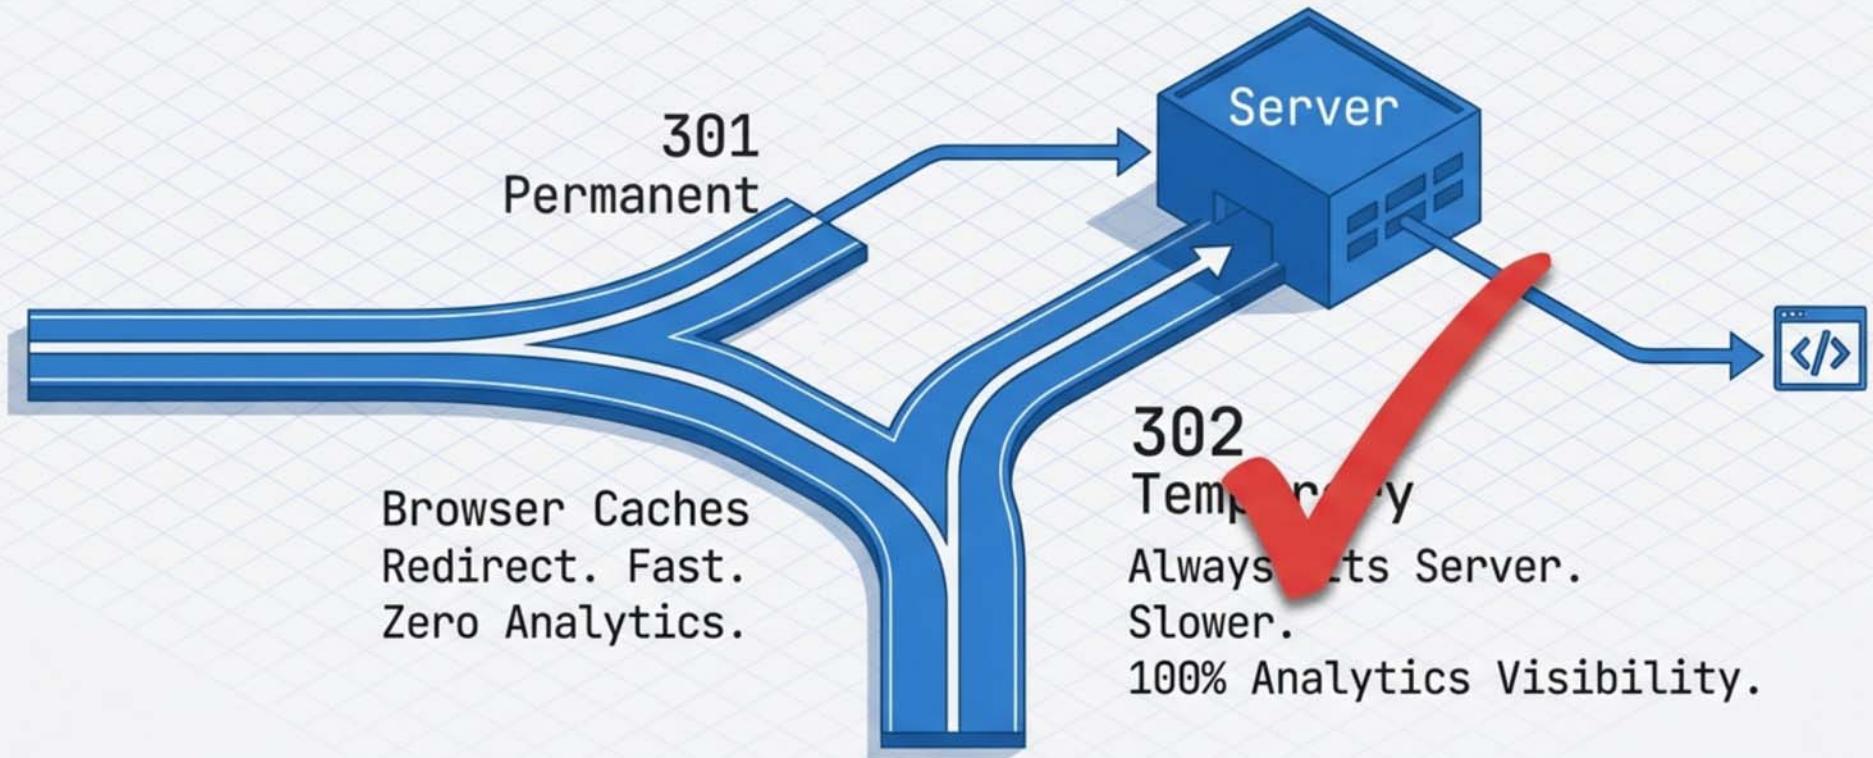
Redirect User

```
GET /api/urls/{shortCode}  
Response: 302 Found  
Location: https://www.example.com/...
```

Phase 4: High-Level Design (The Happy Path)



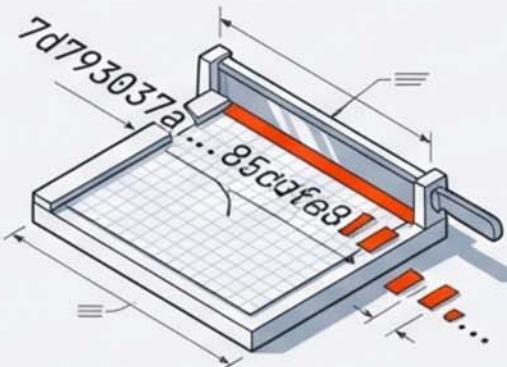
The Redirect Decision: 301 vs. 302



Business Requirement: Analytics drive revenue.

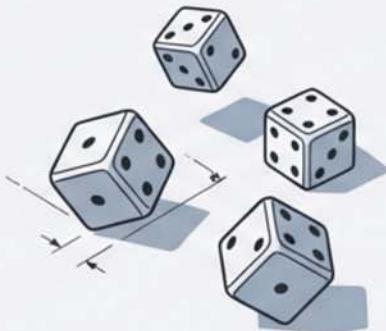
Generating the Short Code

MD5 Hashing



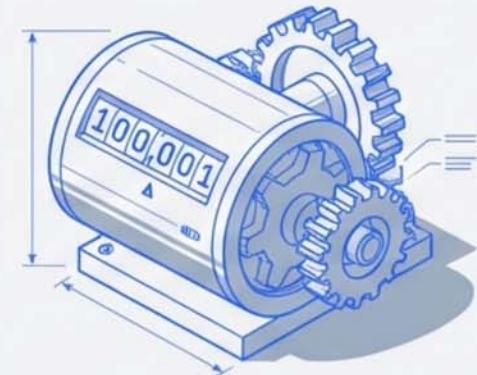
Too long.
Collisions when
truncated.

Random Gen



Collision risk
increases
over time. **Requires**
"Check-Then-Write".

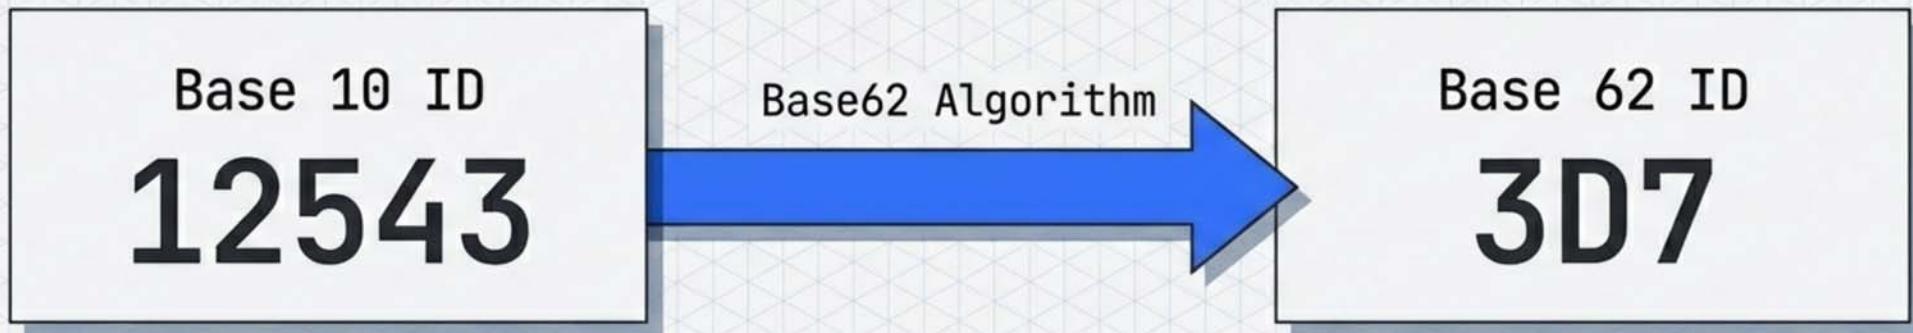
The Counter



Zero Collisions.
Sequential.
Fast.

Modern Architectural Schematic

Base62 Encoding

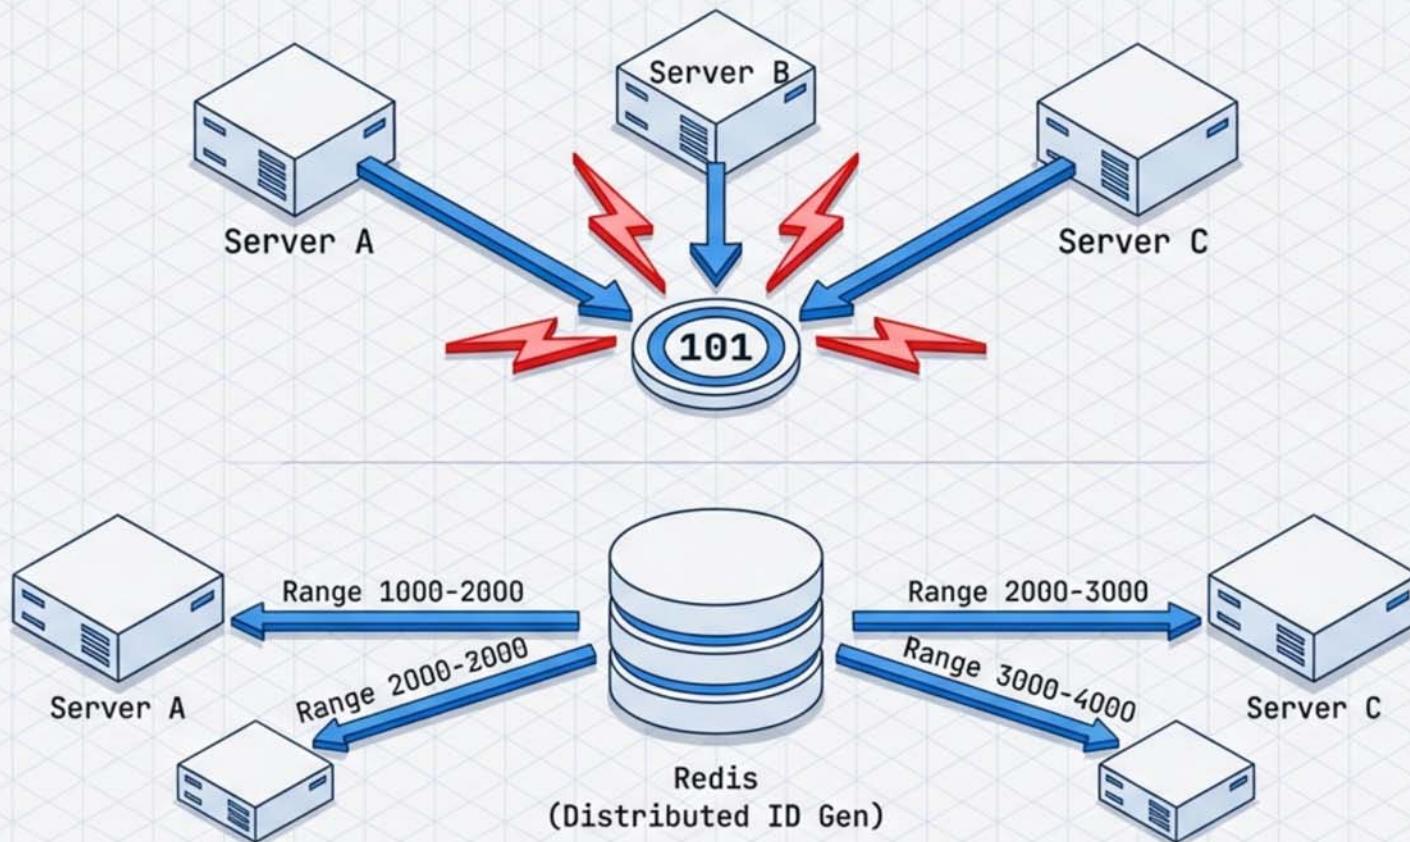


$$[a-z] (26) + [A-Z] (26) + [0-9] (10) = 62 \text{ Characters}$$

\$62^6 \approx 56.8 \text{ Billion Combinations}

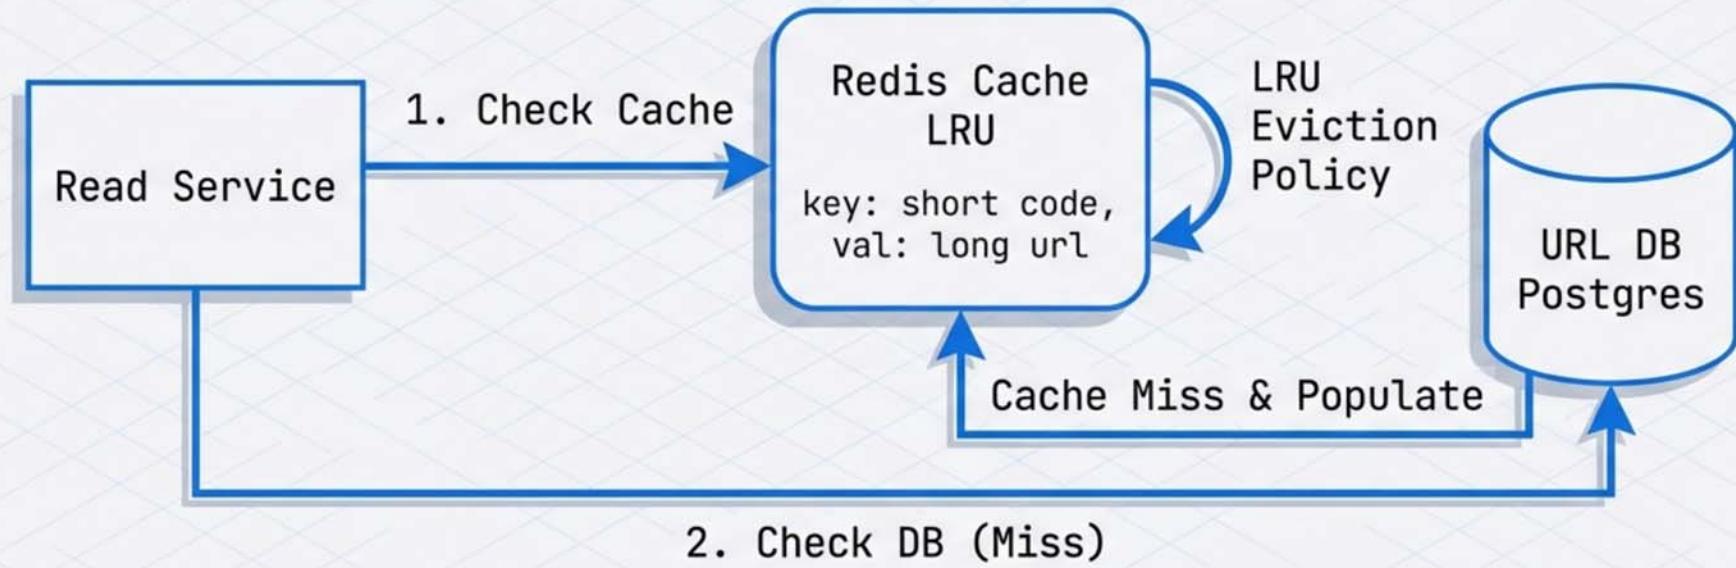
6 characters are enough for 56 Billion URLs.

The Concurrency Problem



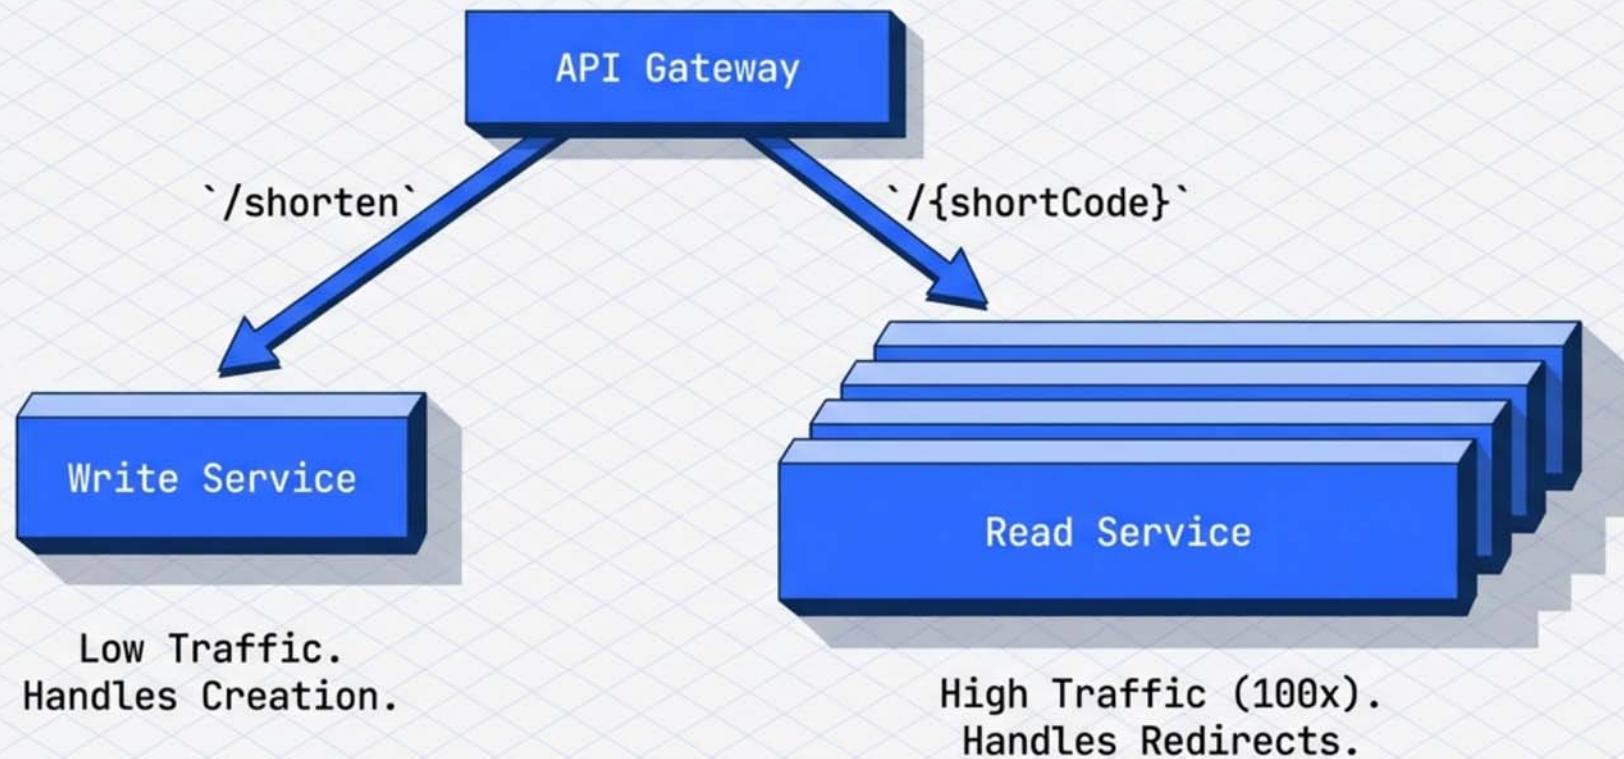
Atomic Increment + Range Batching = No Collisions.

Scaling Reads: The Cache Layer



Handles 99% of Read Traffic.

Microservices & Horizontal Scaling



Global Scale: The CDN Layer



Reduces latency to <50ms by serving redirects from the Edge.

The Final Architecture

